AI Vibrancy EDA

Taylor Boeckman

4/22/2020

```
aiVibrancyIndicators2019 = read.csv("2019indicators.csv")

aiVibrancyIndicators2019.df <- aiVibrancyIndicators2019

aiVibrancyIndicators2019.tib <- as_tibble(aiVibrancyIndicators2019.df) %>%
  mutate( female_ai_authors = as.double( female_ai_authors ) ) %>%
  mutate( female_ai_skill_penetration = as.double( female_ai_skill_penetration ) )
```

```
aiVibrancyIndicators2019.tib %>%
  count( income.group ) %>%
  mutate( pct = n / sum(n),
          pctlabel = paste0( round( pct*100 ), "%" ) ) %>%
  ggplot( aes( x = reorder( income.group, -pct ),
               y = pct ) ) +
      geom_col( fill = "indianred3",
                color = "black" ) +
      geom_text( aes( label = pctlabel ),
                      vjust = -0.25 ) +
      scale_y_continuous( labels = scales::percent ) +
      labs( x = "Income Group",
            y = "Percent",
            title = "Countries Included on Stanford AI Vibrancy Index By Income Group" )
```

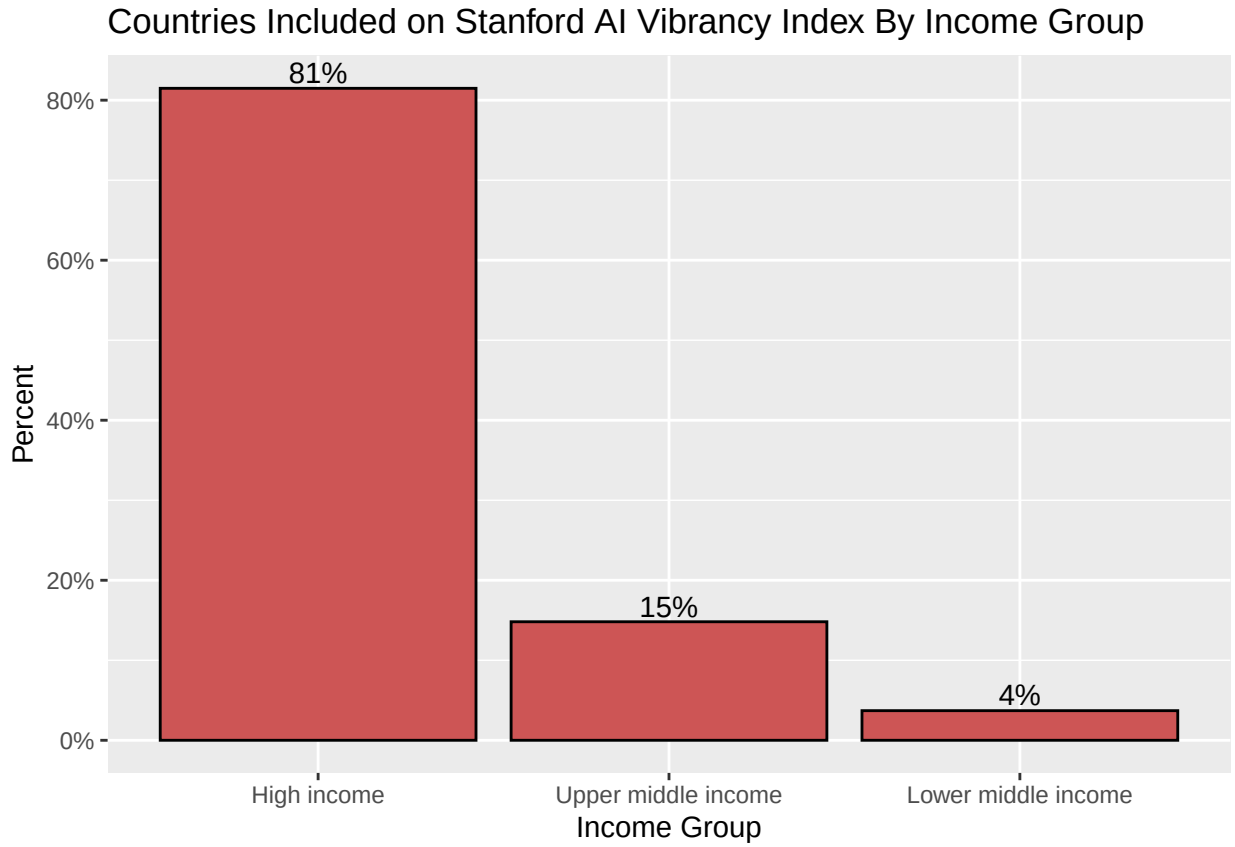


```
aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = ai_privateinvestment_tot , y = country ) ) +
      geom_col( fill = "indianred3",
                color = "black" ) +
      facet_wrap( ~ year, ncol = 2) +
      labs( x = "Private Investment in AI",
            y = "Country",
            title = "Private Investment in AI by Country" )
```

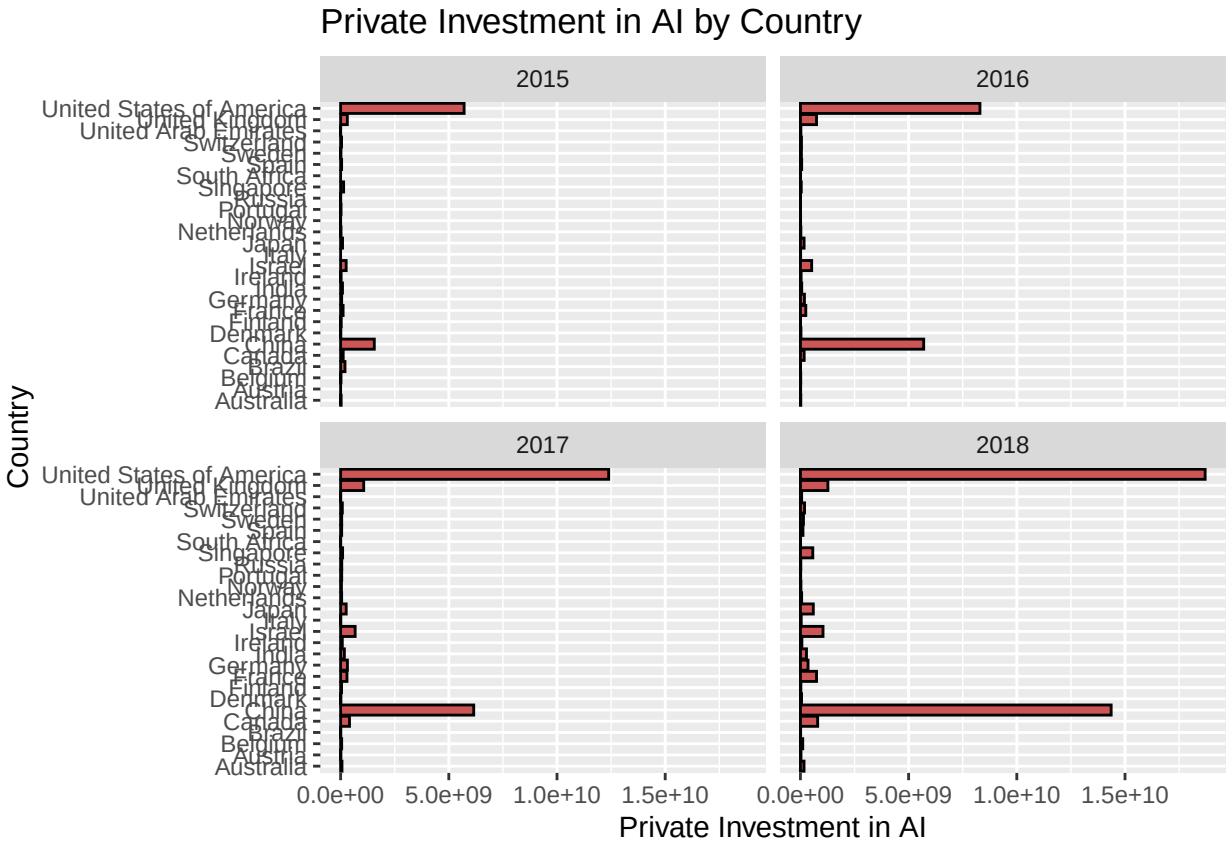


```
aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = ai_privateinvestment_tot_pc , y = country ) ) +
      geom_col( fill = "indianred3",
                color = "black" ) +
      facet_wrap( ~ year, ncol = 2) +
      labs( x = "Private Investment in AI Per Capita",
            y = "Country",
            title = "Private Investment in AI Per Capita by Country" )
```

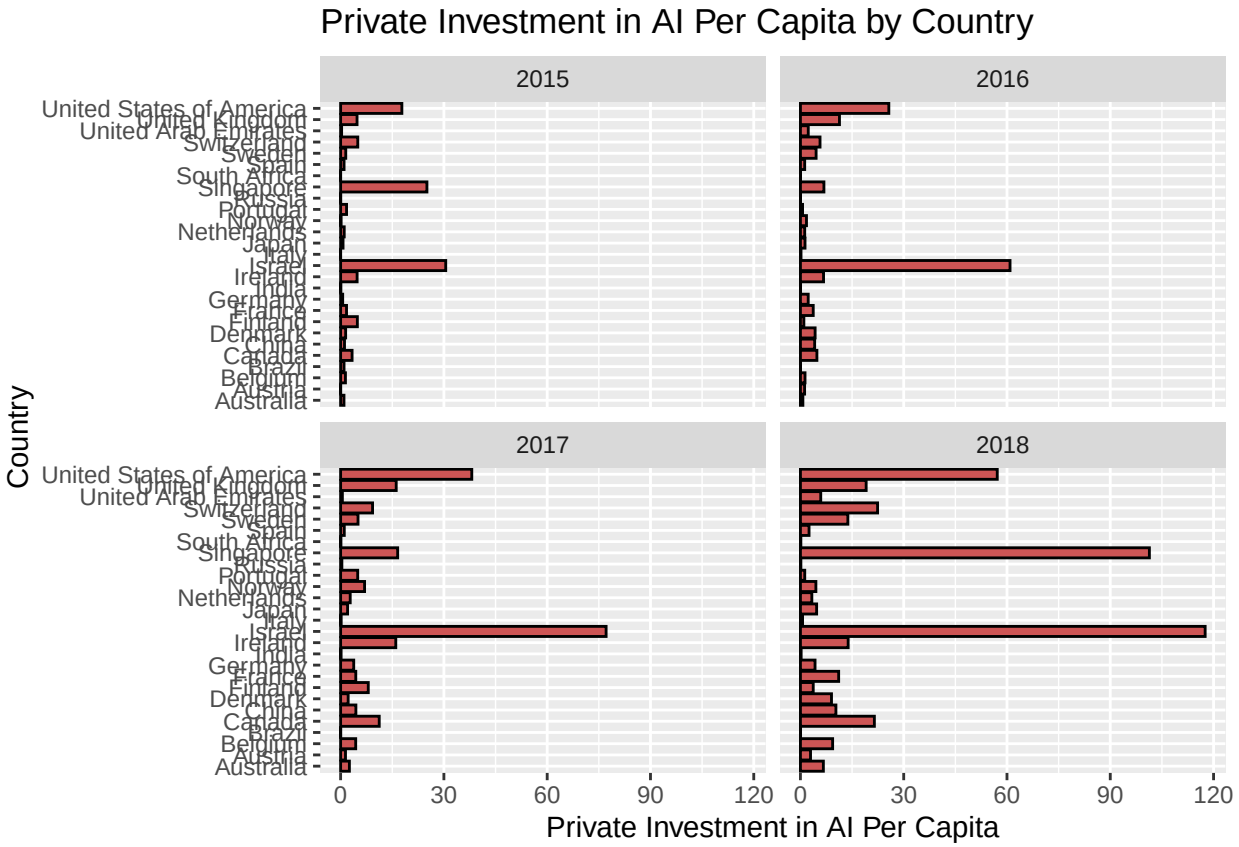


```
num_AIconf_papers_graph <- aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = num_AIconf_papers, y = country, fill = income.group ) ) +
      geom_col() +
      labs( x = "Number of AI Conference Papers",
            y = "Country",
            title = "Number of AI Conference Papers by Country" )

num_AIconf_papers_pc_graph <- aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = num_AIconf_papers_pc, y = country, fill = income.group ) ) +
      geom_col() +
      labs( x = "Number of AI Conference Papers Per Capita",
            y = "Country",
            title = "Number of AI Conference Papers  Per Capita by Country" )

grid.arrange(num_AIconf_papers_graph, num_AIconf_papers_pc_graph, ncol = 2)
```

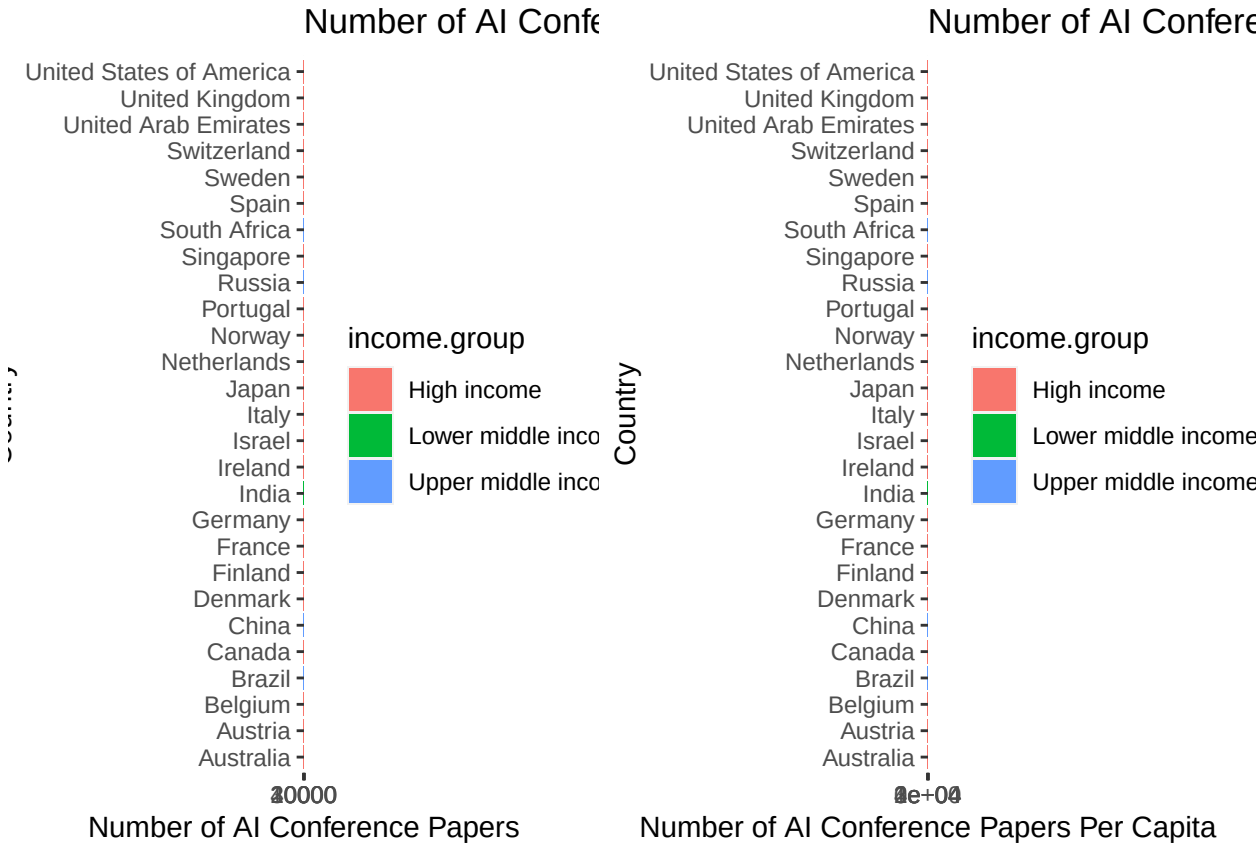


```
aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = num_AIconf_papers, y = country, fill = income.group ) ) +
      geom_col() +
      facet_wrap( ~ year, ncol = 2 ) +
      labs( x = "Number of AI Conference Papers",
            y = "Country",
            title = "Number of AI Conference Papers by Country by Year" )
```

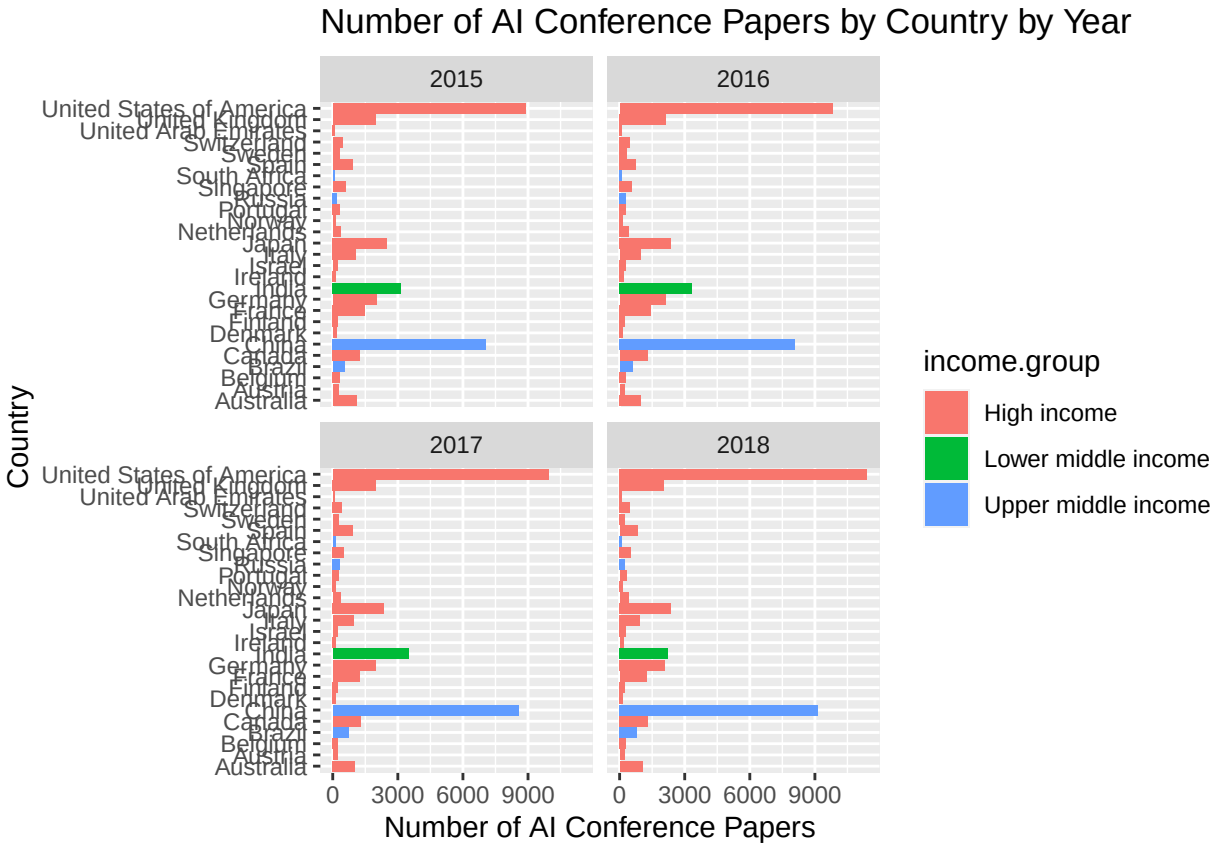


```
aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = num_AIconf_papers_pc, y = country, fill = income.group ) ) +
      geom_col() +
      facet_wrap( ~ year, ncol = 2 ) +
      labs( x = "Number of AI Conference Papers Per Capita",
            y = "Country",
            title = "Number of AI Conference Papers Per Capita by Country by Year" )
```

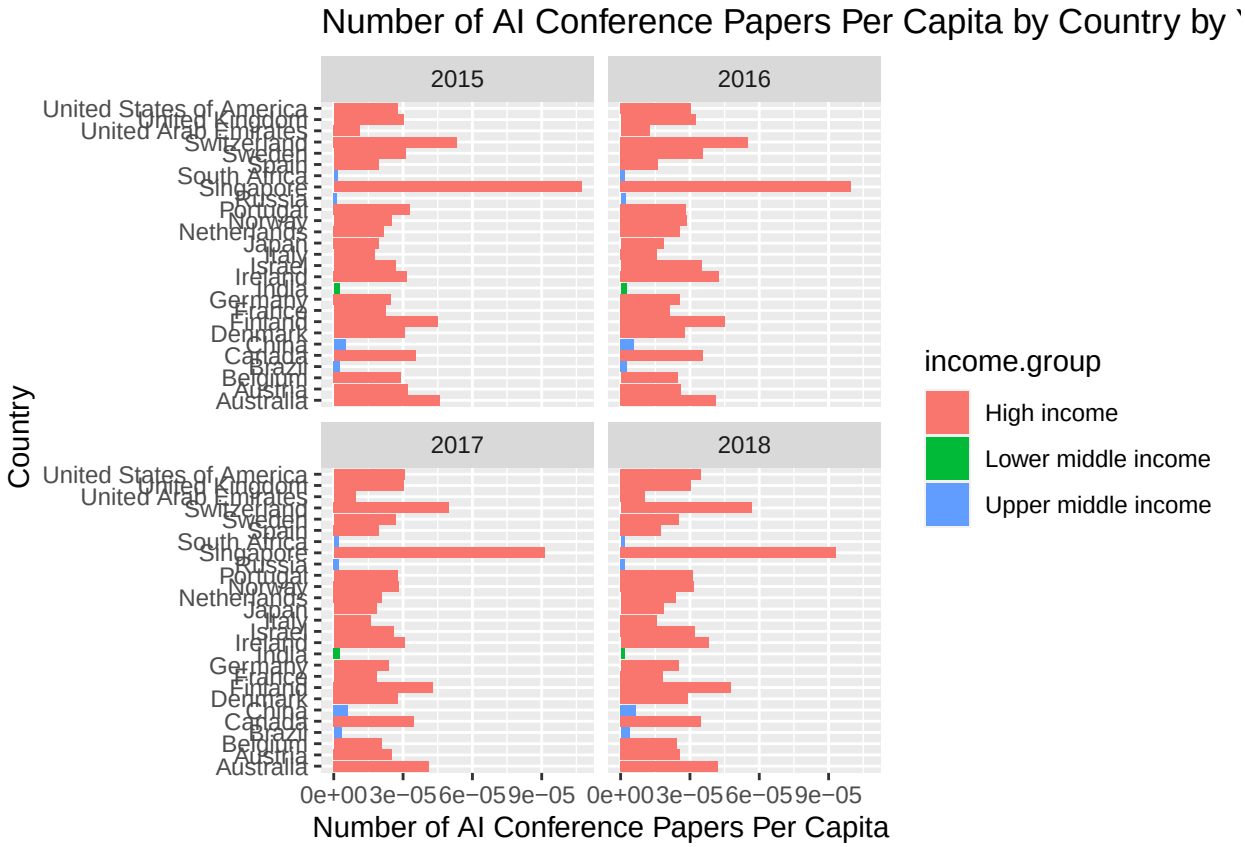


```
num_AIconf_citation_graph <- aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = num_AIconf_citation, y = country ) ) +
      geom_col() +
      labs( x = "Number of AI Conference Citations",
            y = "Country",
            title = "Number of AI Conference Citations by Country" )


num_AIconf_citation_pc_graph <- aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = num_AIconf_citation_pc, y = country ) ) +
      geom_col() +
      labs( x = "Number of AI Conference Citations Per Capita",
            y = "Country",
            title = "Number of AI Conference Citations Per Capita by Country" )

grid.arrange(num_AIconf_citation_graph, num_AIconf_citation_pc_graph, ncol = 2)
```

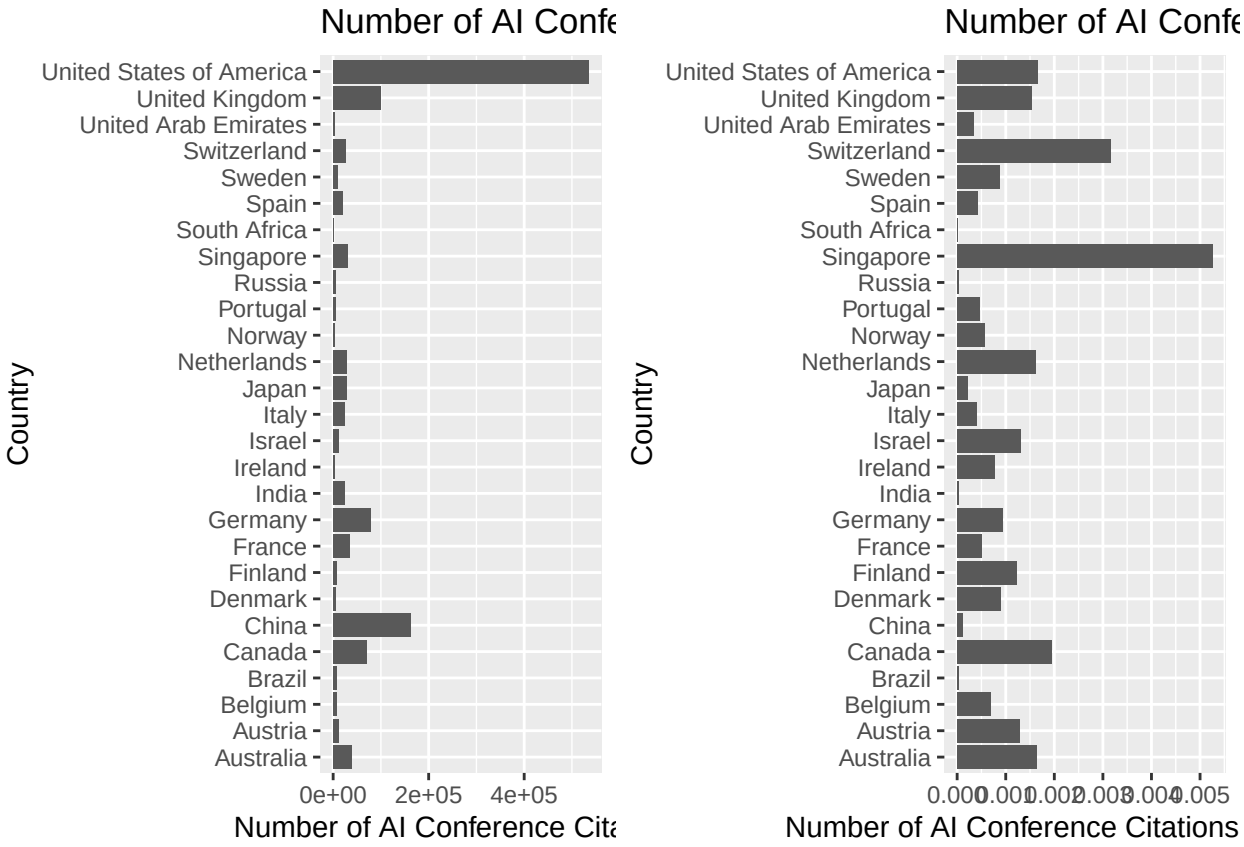


```
aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = num_AIconf_citation, y = country ) ) +
      geom_col() +
      facet_wrap( ~ year, ncol = 2 ) +
      labs( x = "Number of AI Conference Citations",
            y = "Country",
            title = "Number of AI Conference Citations by Country by Year" )
```

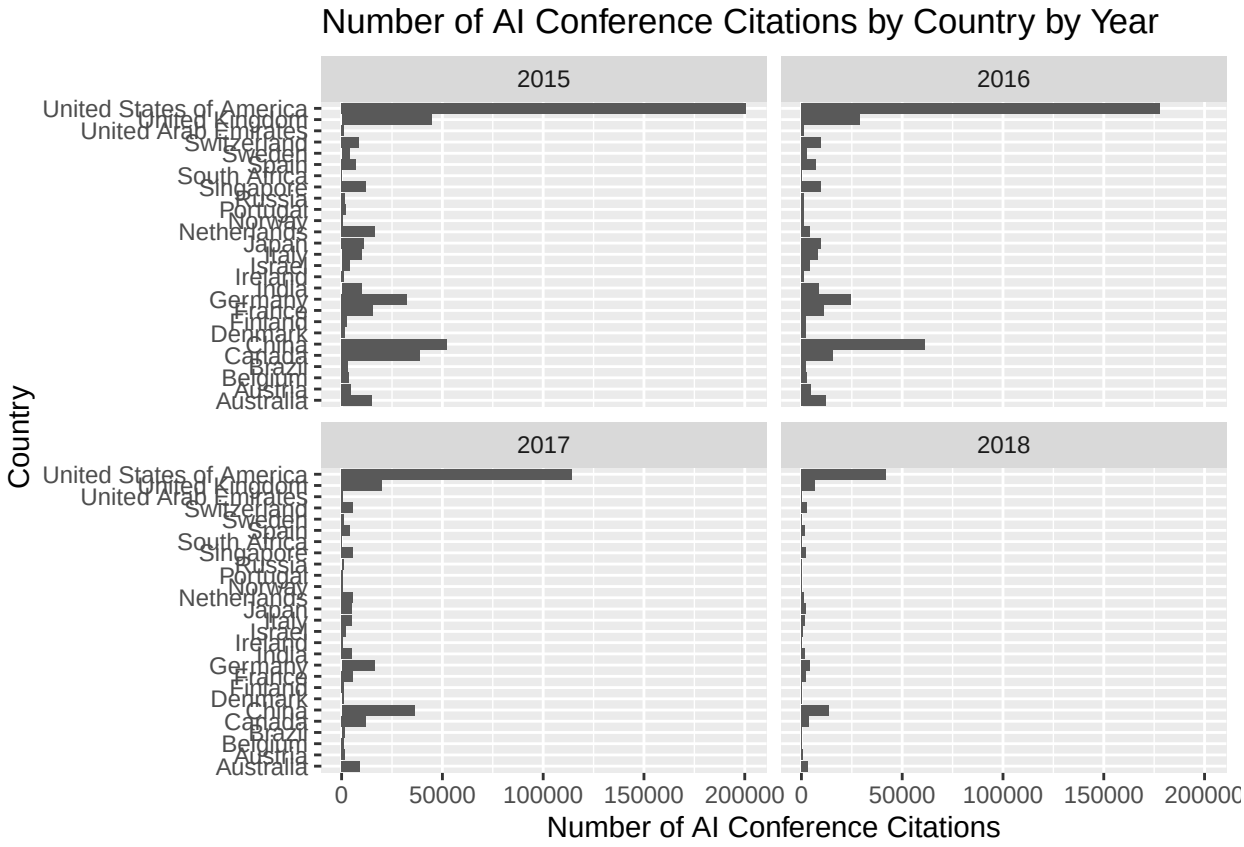


```
aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = num_AIconf_citation_pc, y = country ) ) +
      geom_col() +
      facet_wrap( ~ year, ncol = 2 ) +
      labs( x = "Number of AI Conference Citations Per Capita",
            y = "Country",
            title = "Number of AI Conference Citations Per Capita by Country by Year" )
```

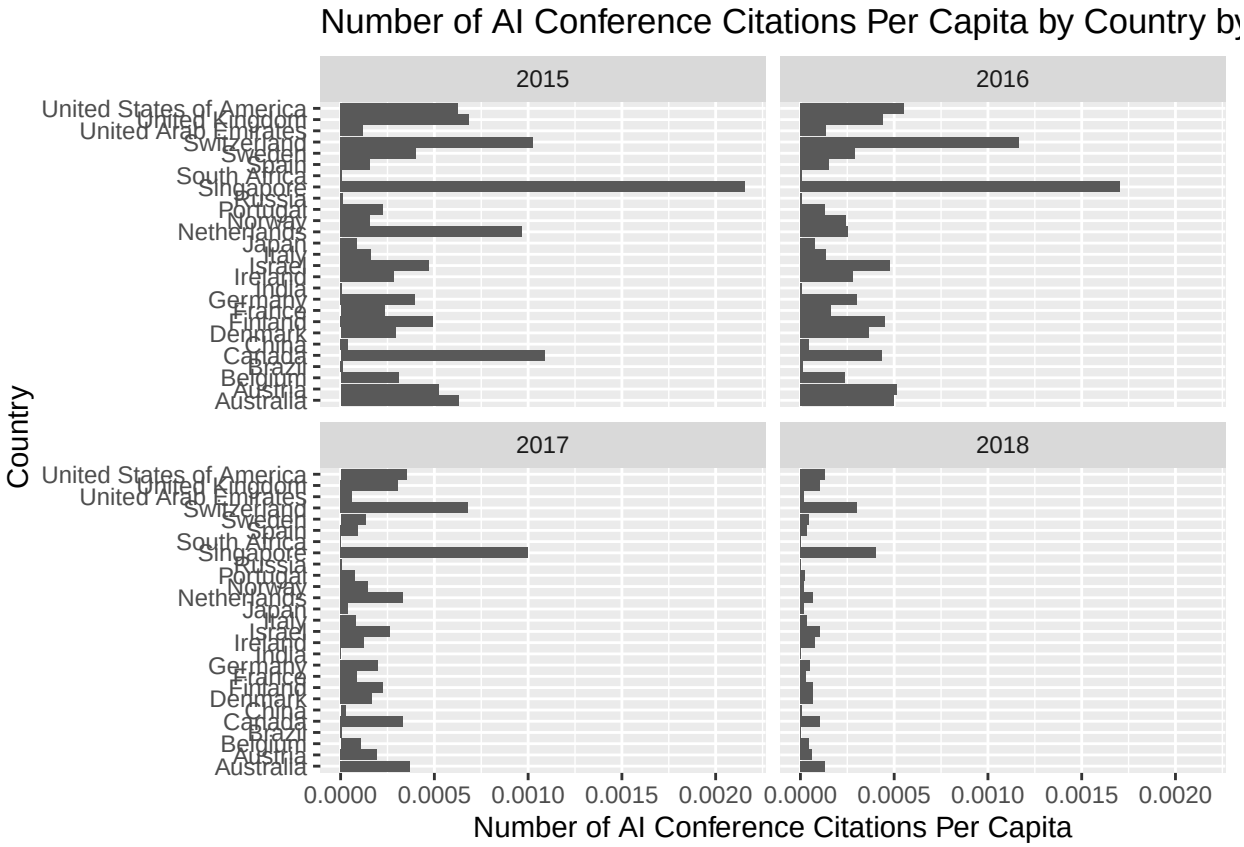


```
aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = country , y = female_ai_authors ) ) +
  geom_col( fill = "indianred3",
                color = "black" ) +
  theme(axis.text.x = element_text(angle = 90))
```

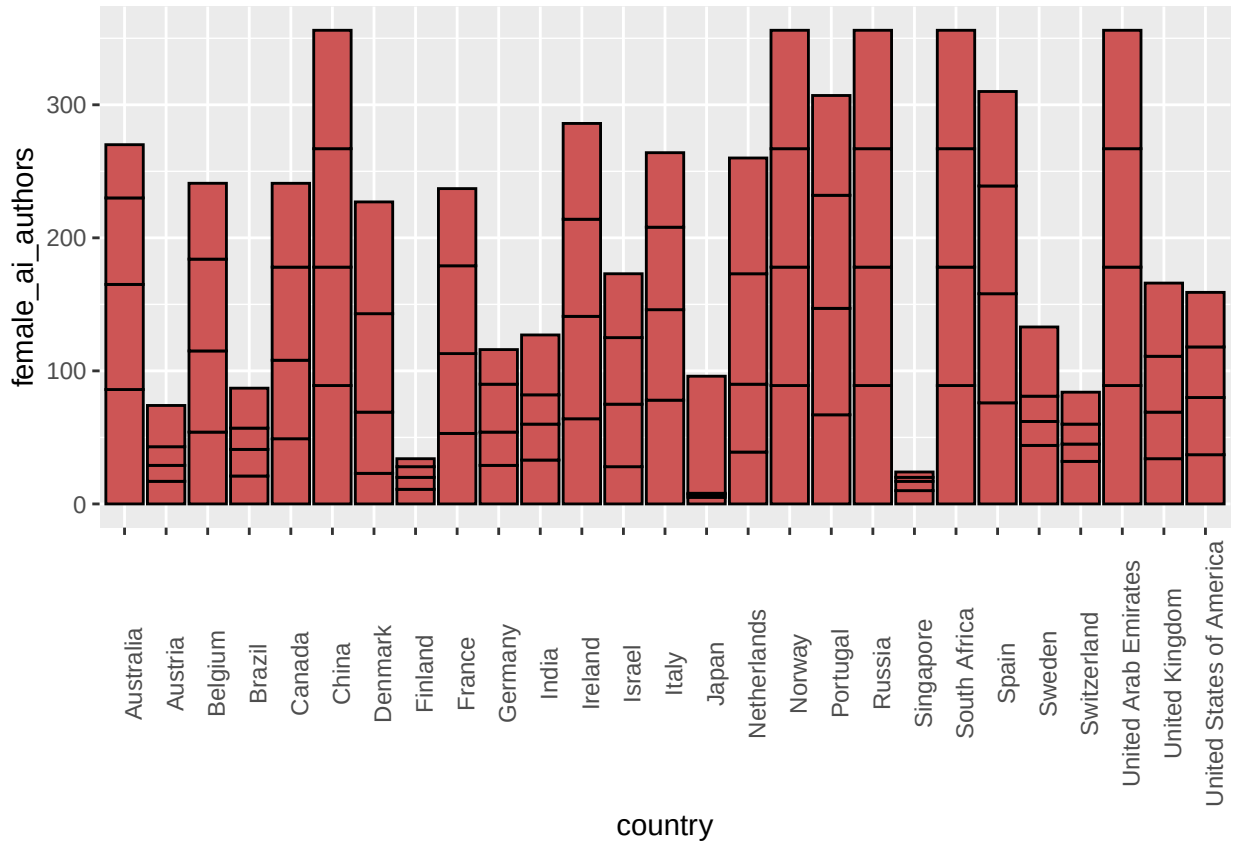


```
aiVibrancyIndicators2019.tib %>%
  filter(year == "2018") %>%
  ggplot( aes( x = country , y = female_ai_skill_penetration) ) +
  geom_col( fill = "indianred3",
                color = "black" ) +
  theme(axis.text.x = element_text(angle = 90))
```

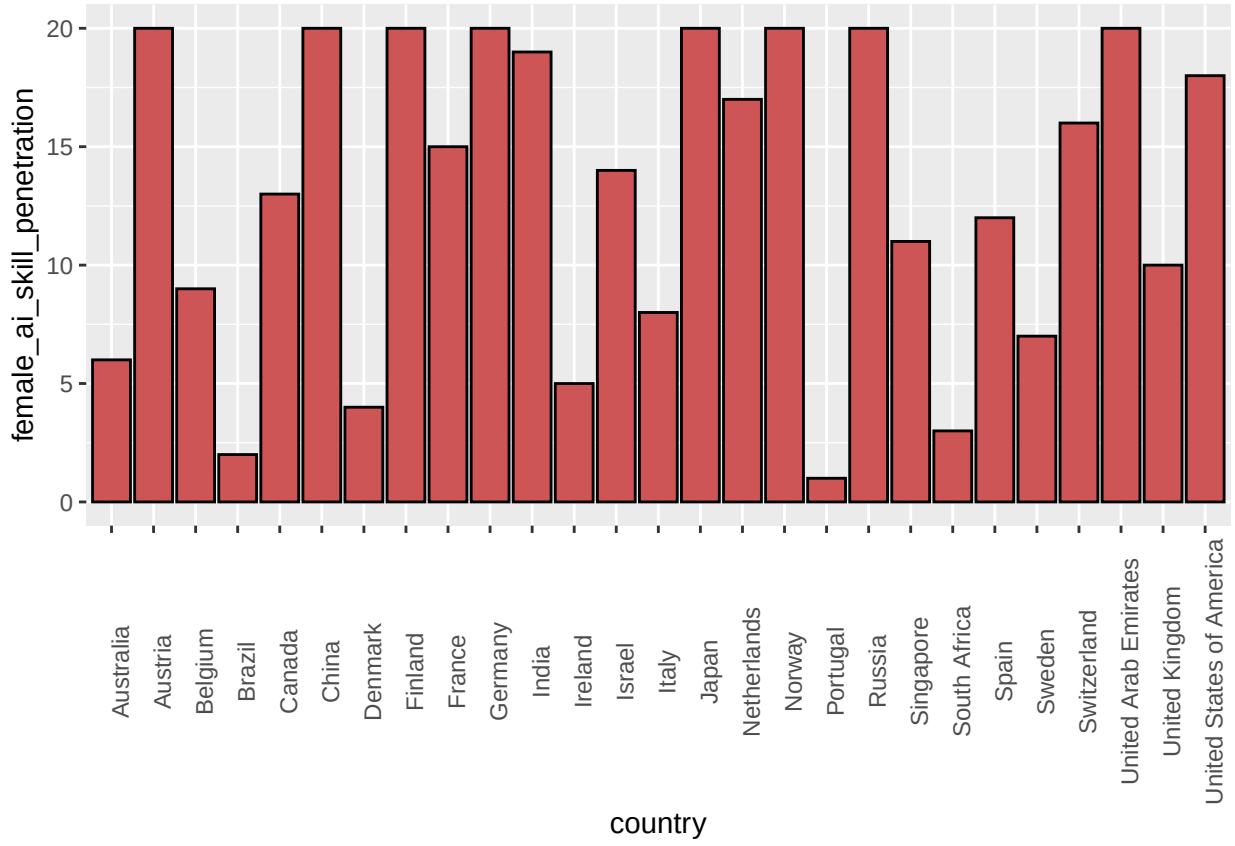


```
aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = country , y = num_AIstartups ) ) +
  geom_col( fill = "indianred3",
                color = "black" ) +
  theme(axis.text.x = element_text(angle = 90))
```

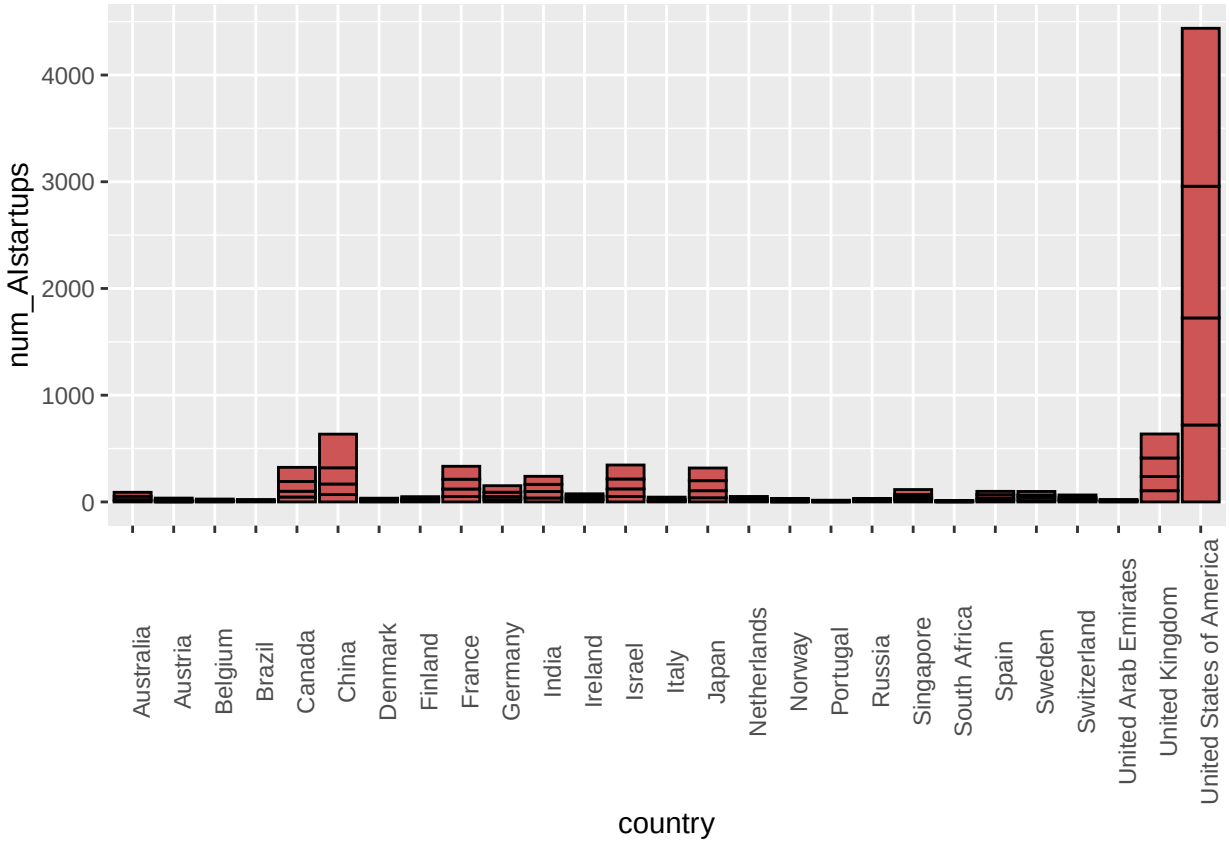


```
aiVibrancyIndicators2019.tib %>%
  ggplot( aes( x = country , y = num_AIstartups_pc ) ) +
  geom_col( fill = "indianred3",
                color = "black" ) +
  theme(axis.text.x = element_text(angle = 90))
```

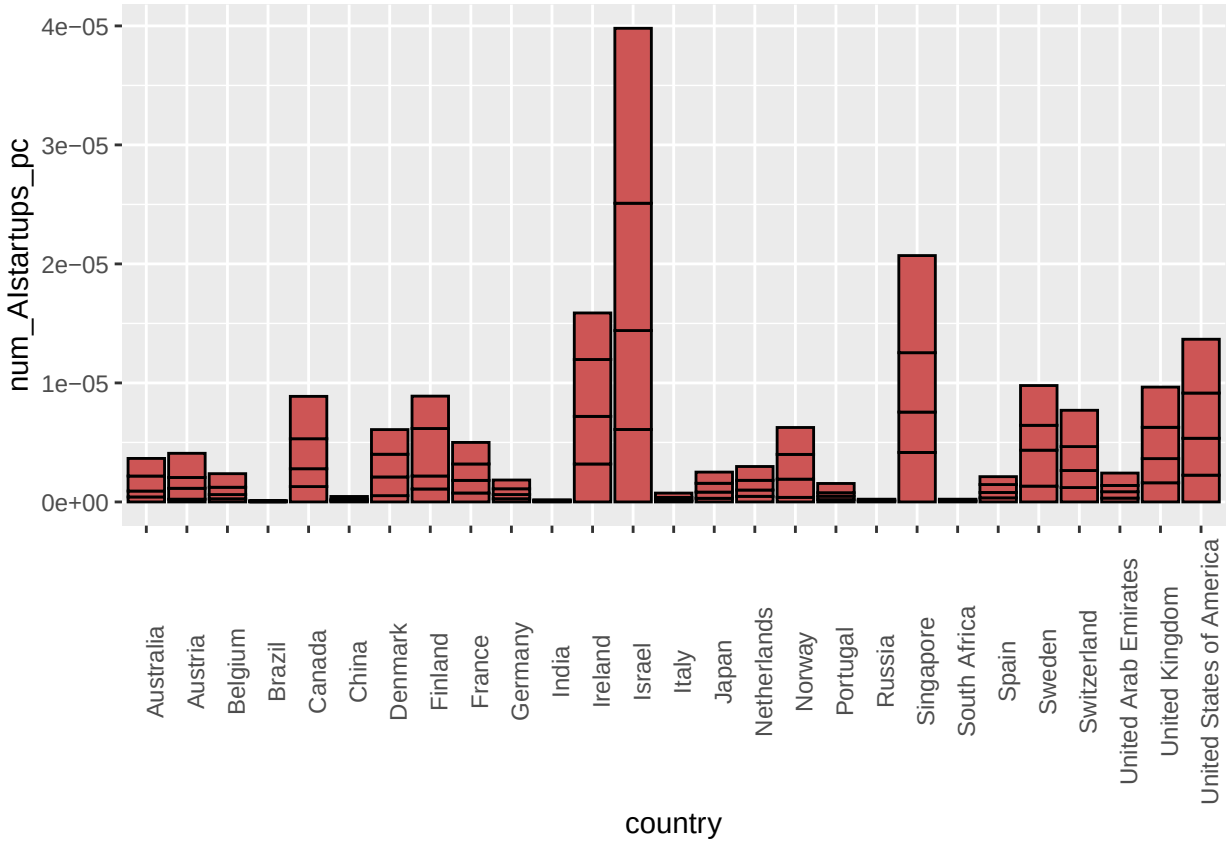


```
aiVibrancyNumeric.tib <- aiVibrancyIndicators2019.tib %>%
  select_if( is.numeric )

r.mat = cor( aiVibrancyNumeric.tib, use = "complete.obs" )

ggcorrplot( r.mat,
            type = "lower",
            insig = "blank" )
```

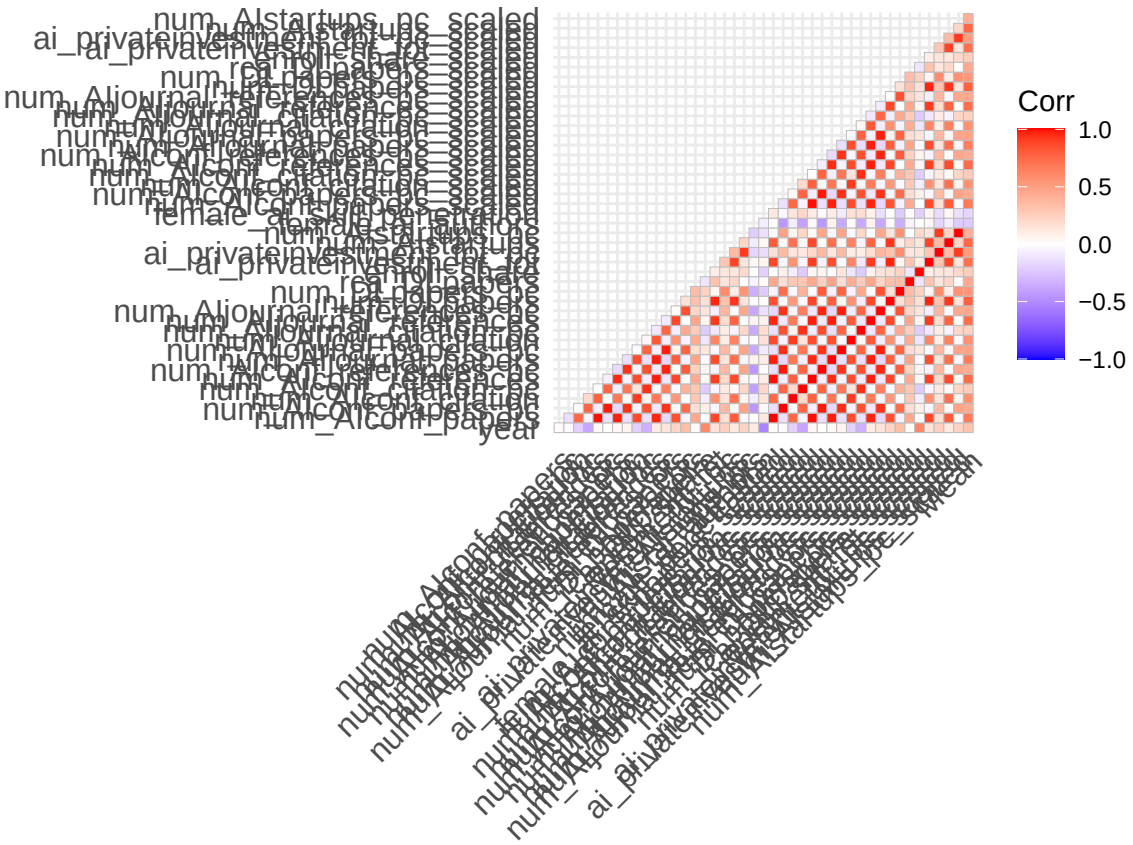


```
gower_dist <- daisy( aiVibrancyIndicators2019.tib %>% select( -country ),
                     metric = "gower",
                     type = list( logratio = 3 )
                   )

gower.summ <- summary( gower_dist )
gower_mat <- as.matrix(gower_dist)

closest.tib <-
aiVibrancyIndicators2019.tib[
        which( gower_mat == min(gower_mat[ gower_mat != min(gower_mat)]),
               arr.ind = TRUE)[1, ], ]

farthest.tib <-
aiVibrancyIndicators2019.tib[
      which(gower_mat == max(gower_mat[gower_mat != max(gower_mat)]),
              arr.ind = TRUE)[1, ], ]

k.vec <- 1:15

get_pam_silwidth <- function( k, dist )
{
  pam.clust <- pam( dist, diss=TRUE, k=k )
  return( pam.clust$silinfo$avg.width )
}
```

```
sil_width.vec <- map_dbl( k.vec[ -1 ],
  get_pam_silwidth,
dist=gower_dist
  )

pam_sil.tib <- tibble( k = k.vec,
  sil = c( 0, sil_width.vec )
)

sil_max <- with( pam_sil.tib,
  which( sil == max( sil ) )
)

pam.clust <- pam( x = gower_dist,
  k = sil_max,
  diss = TRUE
)

aiVibrancyIndicators2019.tib <- aiVibrancyIndicators2019.tib %>%
    mutate(p3 = factor( paste0( 'p', pam.clust$clustering ) ) )

aiVibrancyIndicators2019_vars.vec <-
  c("country", "num_AIconf_papers","num_AIconf_papers_pc", "num_AIconf_citation", "num_AIconf_citation_p

aiVibrancyIndicators2019.famd <-
  FAMD( aiVibrancyIndicators2019.tib %>%
          select(all_of(aiVibrancyIndicators2019_vars.vec)), ncp = 6, graph = FALSE )

( aiVibrancyIndicators2019.scree.ggplot <-
  fviz_screeplot( aiVibrancyIndicators2019.famd ) %>%
  labs( title = "Scree Plot") )
```

```
## [[1]]
```

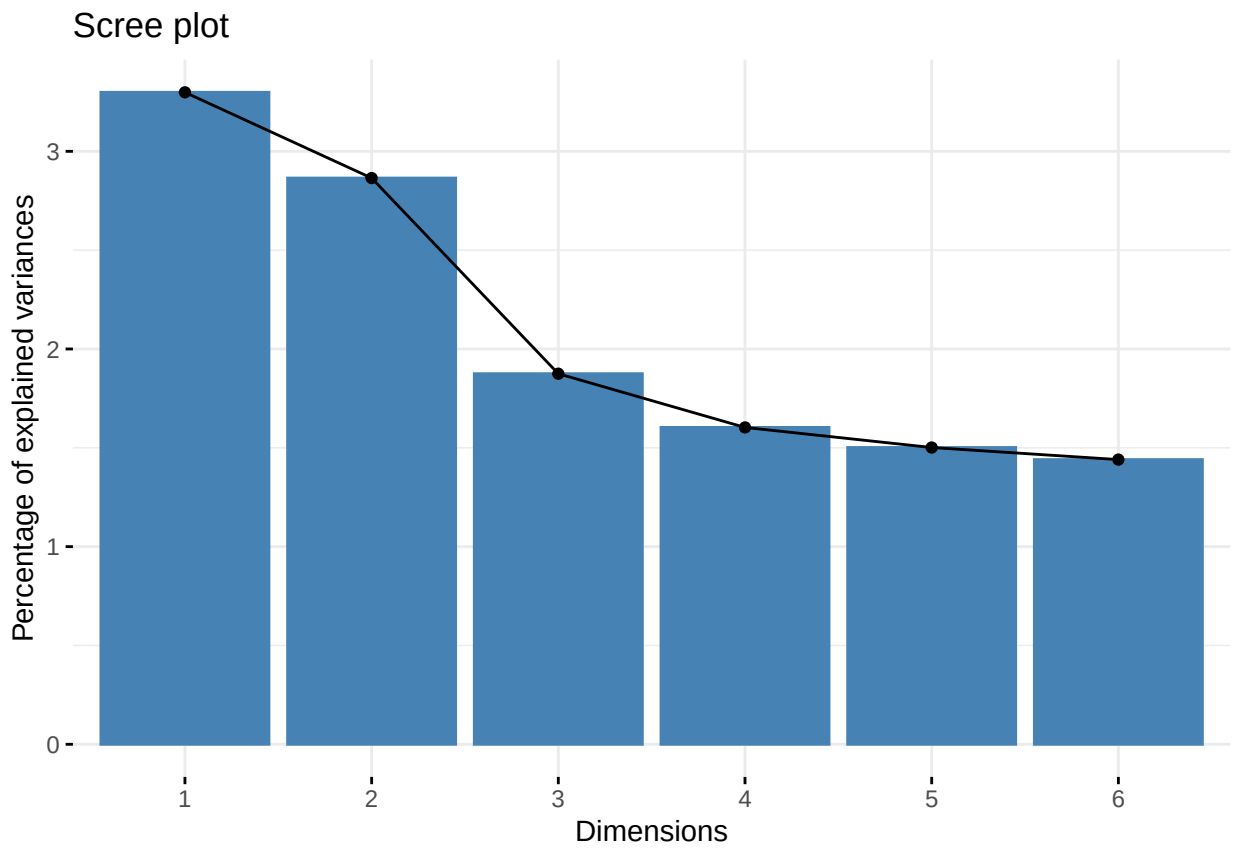


```
##
## $title
## [1] "Scree Plot"
##
## attr(,"class")
## [1] "labels"
```

```
( aiVibrancyIndicators2019.var.ggplot <-
      fviz_famd_var( aiVibrancyIndicators2019.famd, repel = TRUE ) +
      labs(title = "FAMD Variables Contributions to the Dimensions"))
```

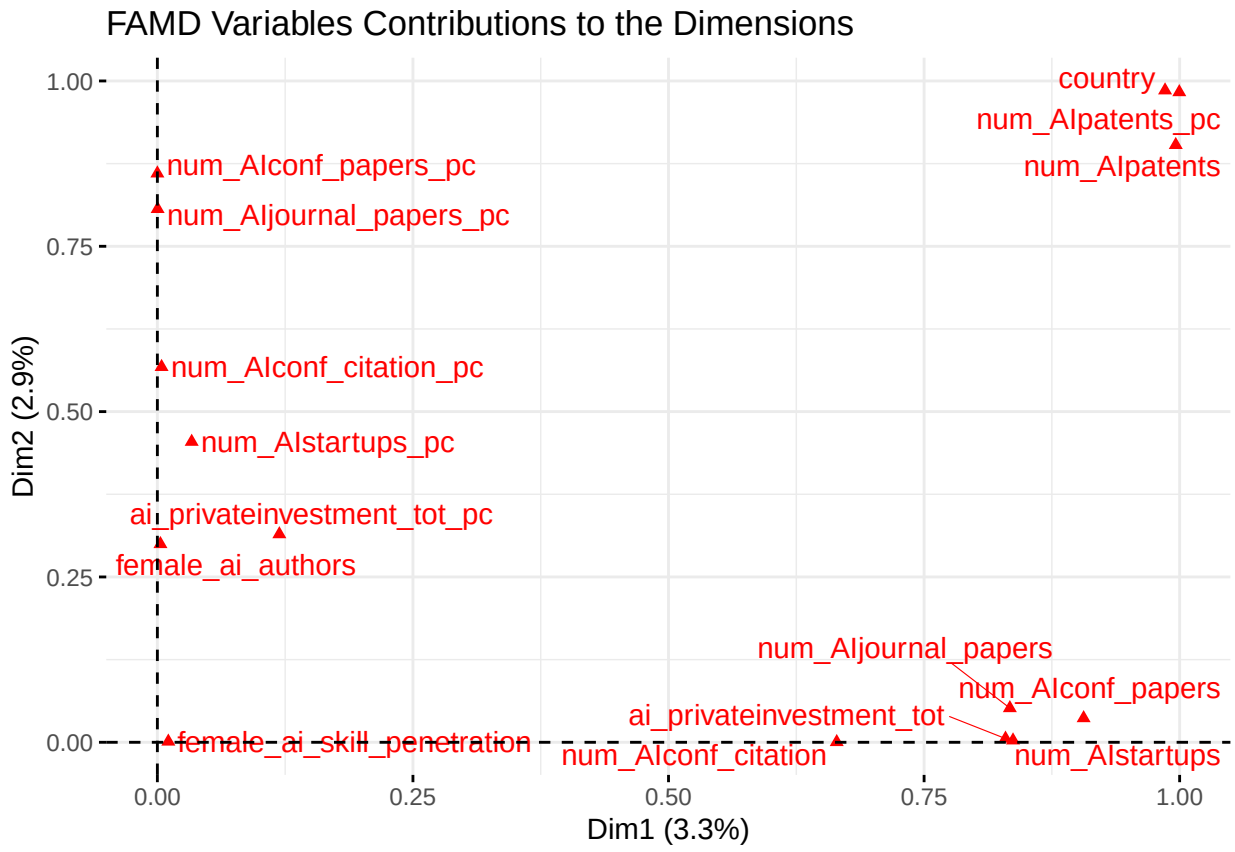


```
aiVibrancyIndicators2019.contrib.dim1.ggplot <-
      fviz_contrib( aiVibrancyIndicators2019.famd, "var", axes = 1,
                    title = "Contribution of vars to Dim-1" )

aiVibrancyIndicators2019.contrib.dim2.ggplot <-
      fviz_contrib( aiVibrancyIndicators2019.famd, "var", axes = 2,
                    title = "Contribution of vars to Dim-2" )

contrib.grid <-
    grid.arrange(aiVibrancyIndicators2019.contrib.dim1.ggplot,
                 aiVibrancyIndicators2019.contrib.dim2.ggplot, nrow = 1,
                top = "Sampled IPs Quality of Representation")
```

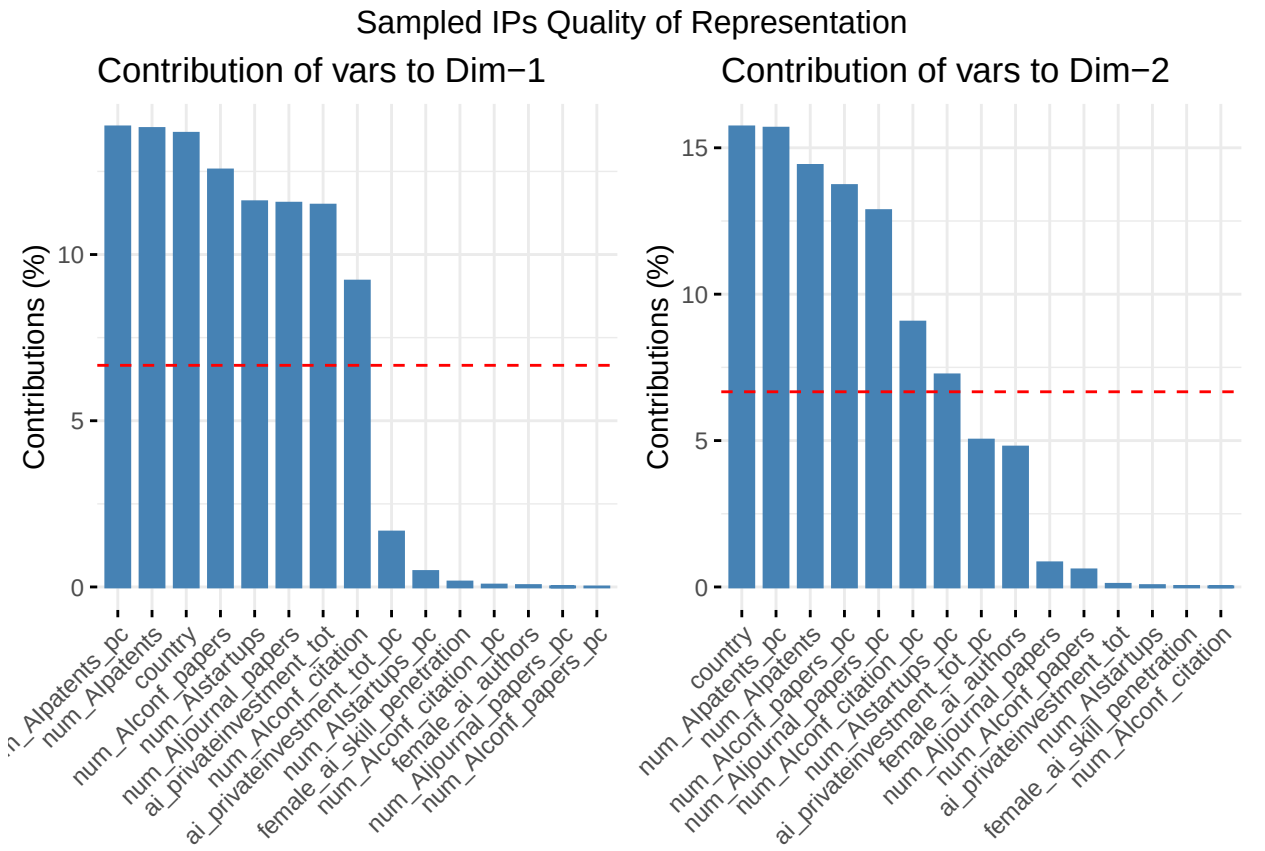


```
( pam_sil.ggplot <- pam_sil.tib %>%
  ggplot( mapping = aes( x = k, y = sil ) ) +
              geom_line() +
              geom_vline( xintercept = sil_max,
              linetype="dashed", color="blue" ) +
              theme_minimal() +
              scale_x_continuous( breaks = k.vec ) +
              labs( title = "PAM Silhouette Widths" ) +
              ylab( "silhouette width" ) )
```

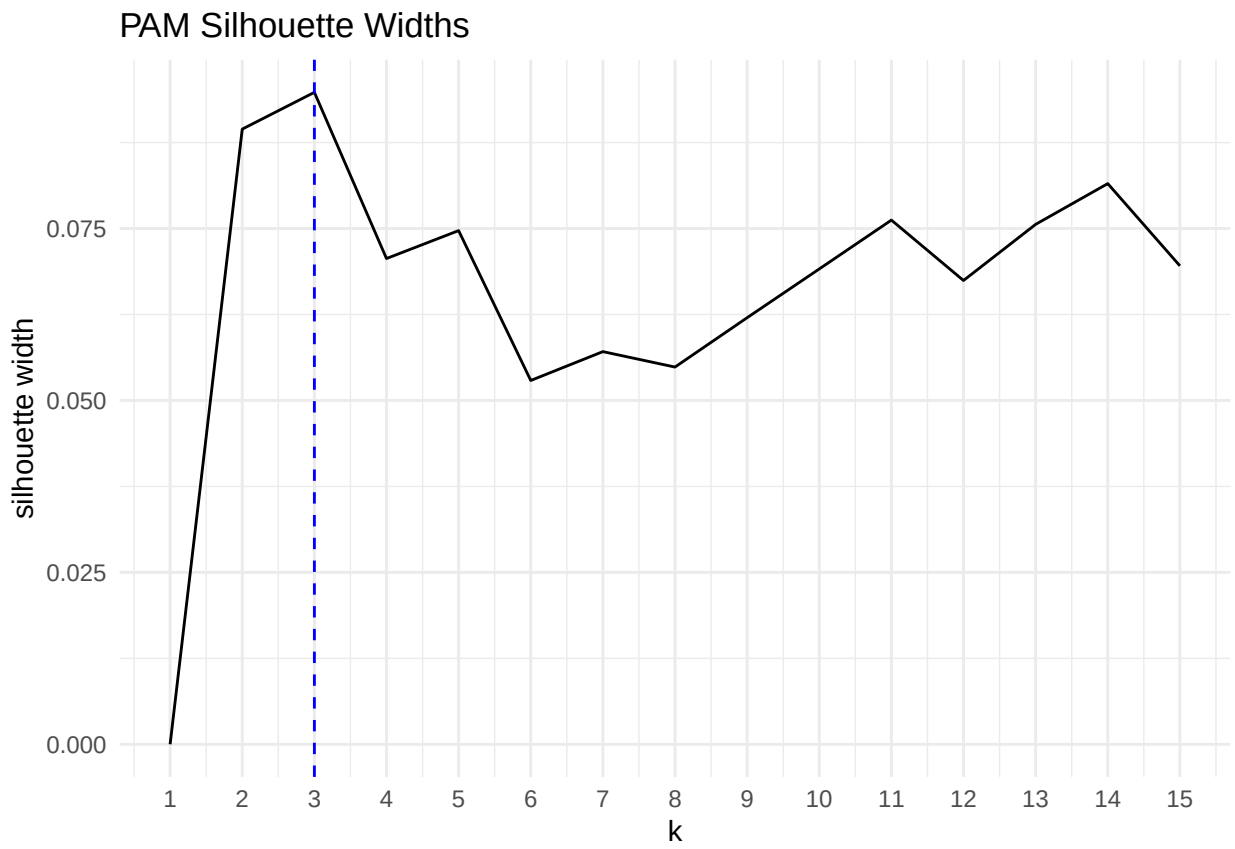


```
aiVibrancyIndicators2019.tib <- aiVibrancyIndicators2019.tib %>%
  mutate( famd_dim1 = aiVibrancyIndicators2019.famd$ind$coord[, 1],
          famd_dim2 = aiVibrancyIndicators2019.famd$ind$coord[, 2]
        )

pam_clusters_guide = "PAM\nClusters"

pam_cluster_colors.vec <- brewer.pal( sil_max, name="Set1" )
names(pam_cluster_colors.vec) <- paste0( 'p', 1:sil_max )

quali_ind.tib <- tibble( var  = rownames( aiVibrancyIndicators2019.famd$quali.var$coord ),
                         dim1 = aiVibrancyIndicators2019.famd$quali.var$coord[, 1],
                         dim2 = aiVibrancyIndicators2019.famd$quali.var$coord[, 2]
                       )

(aiVibrancyIndicators2019.ggplot <- aiVibrancyIndicators2019.tib %>%
  ggplot( mapping = aes( x = famd_dim1, y = famd_dim2) ) +
    geom_vline( xintercept = 0 ) +
    geom_hline( yintercept = 0 ) +
    geom_point( mapping = aes( color = p3 ),
                alpha = .5 ) +
    geom_text(aes(label= country),hjust=0, vjust=0, alpha = .7) +
    geom_encircle( mapping = aes( group = p3, color = p3 ),
                   linetype = "dotted", s_shape = 0.95 ) +
    theme_minimal() +
    coord_cartesian( xlim = c(-2.5, 20), ylim = c(-5, 15) ) +
```

```
    labs( title = "PAM Clusters + FAMD Decomposition" ) )
```